

In this discussion, we will build a basic hybrid CNN-LSTM model for classification on the EEG dataset

This notebook was inspired to Tonmoy with some attempts to tune by us

(i) Importing the necessary packages

```
In [ ]: import numpy as np
import pandas as pd
import keras
from keras.models import Sequential, Model
from keras.layers import Dense, Activation, Flatten, Dropout
from keras.layers import Conv2D, LSTM, BatchNormalization, MaxPooling2D, Reshape
from keras.utils import to_categorical
import matplotlib.pyplot as plt
```

```
2023-03-18 20:47:02.872775: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
```

(ii) Preprocessing the dataset and preparing the training, validation, and test datasets

```

In [ ]: def data_prep(X,y,sub_sample,average,noise):

    total_X = None
    total_y = None

    # Trimming the data (sample,22,1000) -> (sample,22,500)
    X = X[:, :, 0:500]
    print('Shape of X after trimming:', X.shape)

    # Maxpooling the data (sample,22,1000) -> (sample,22,500/sub_sample)
    X_max = np.max(X.reshape(X.shape[0], X.shape[1], -1, sub_sample), axis=3)

    total_X = X_max
    total_y = y
    print('Shape of X after maxpooling:', total_X.shape)

    # Averaging + noise
    X_average = np.mean(X.reshape(X.shape[0], X.shape[1], -1, average), axis=3)
    X_average = X_average + np.random.normal(0.0, 0.5, X_average.shape)

    total_X = np.vstack((total_X, X_average))
    total_y = np.hstack((total_y, y))
    print('Shape of X after averaging+noise and concatenating:', total_X.shape)

    # Subsampling
    for i in range(sub_sample):

        X_subsample = X[:, :, i::sub_sample] + \
            (np.random.normal(0.0, 0.5, X[:, :, i::sub_sample].shape))

        total_X = np.vstack((total_X, X_subsample))
        total_y = np.hstack((total_y, y))

    print('Shape of X after subsampling and concatenating:', total_X.shape)
    return total_X, total_y

```

```

In [ ]: ## Loading the dataset

X_test = np.load("X_test.npy")
y_test = np.load("y_test.npy")
person_train_valid = np.load("person_train_valid.npy")
X_train_valid = np.load("X_train_valid.npy")
y_train_valid = np.load("y_train_valid.npy")
person_test = np.load("person_test.npy")

## Adjusting the labels so that

# Cue onset left - 0
# Cue onset right - 1
# Cue onset foot - 2
# Cue onset tongue - 3

y_train_valid -= 769
y_test -= 769

## Random splitting and reshaping the data
# First generating the training and validation indices using random splitting

ind_valid = np.random.choice(2115, 375, replace=False)
ind_train = np.array(list(set(range(2115)).difference(set(ind_valid))))

# Creating the training and validation sets using the generated indices
(X_train, X_valid) = X_train_valid[ind_train], X_train_valid[ind_valid]
(y_train, y_valid) = y_train_valid[ind_train], y_train_valid[ind_valid]

## Preprocessing the dataset
x_train, y_train = data_prep(X_train, y_train, 2, 2, True)
x_valid, y_valid = data_prep(X_valid, y_valid, 2, 2, True)
X_test_prep, y_test_prep = data_prep(X_test, y_test, 2, 2, True)

print('Shape of training set:', x_train.shape)
print('Shape of validation set:', x_valid.shape)
print('Shape of training labels:', y_train.shape)
print('Shape of validation labels:', y_valid.shape)
print('Shape of testing set:', X_test_prep.shape)
print('Shape of testing labels:', y_test_prep.shape)

# Converting the labels to categorical variables for multiclass classification
y_train = to_categorical(y_train, 4)
y_valid = to_categorical(y_valid, 4)
y_test = to_categorical(y_test_prep, 4)
print('Shape of training labels after categorical conversion:', y_train.shape)
print('Shape of validation labels after categorical conversion:', y_valid.shape)
print('Shape of test labels after categorical conversion:', y_test.shape)

# Adding width of the segment to be 1
x_train = x_train.reshape(x_train.shape[0], x_train.shape[1], x_train.shape[2], 1)

```

```

x_valid = x_valid.reshape(x_valid.shape[0], x_valid.shape[1], x_train.shape[1])
x_test = X_test_prep.reshape(X_test_prep.shape[0], X_test_prep.shape[1], X_train.shape[1])
print('Shape of training set after adding width info:',x_train.shape)
print('Shape of validation set after adding width info:',x_valid.shape)
print('Shape of test set after adding width info:',x_test.shape)

# Reshaping the training and validation dataset
x_train = np.swapaxes(x_train, 1,3)
x_train = np.swapaxes(x_train, 1,2)
x_valid = np.swapaxes(x_valid, 1,3)
x_valid = np.swapaxes(x_valid, 1,2)
x_test = np.swapaxes(x_test, 1,3)
x_test = np.swapaxes(x_test, 1,2)
print('Shape of training set after dimension reshaping:',x_train.shape)
print('Shape of validation set after dimension reshaping:',x_valid.shape)
print('Shape of test set after dimension reshaping:',x_test.shape)

keras.backend.clear_session()

Shape of X after trimming: (1740, 22, 500)
Shape of X after maxpooling: (1740, 22, 250)
Shape of X after averaging+noise and concatenating: (3480, 22, 250)
Shape of X after subsampling and concatenating: (6960, 22, 250)
Shape of X after trimming: (375, 22, 500)
Shape of X after maxpooling: (375, 22, 250)
Shape of X after averaging+noise and concatenating: (750, 22, 250)
Shape of X after subsampling and concatenating: (1500, 22, 250)
Shape of X after trimming: (443, 22, 500)
Shape of X after maxpooling: (443, 22, 250)
Shape of X after averaging+noise and concatenating: (886, 22, 250)
Shape of X after subsampling and concatenating: (1772, 22, 250)
Shape of training set: (6960, 22, 250)
Shape of validation set: (1500, 22, 250)
Shape of training labels: (6960,)
Shape of validation labels: (1500,)
Shape of testing set: (1772, 22, 250)
Shape of testing labels: (1772,)
Shape of training labels after categorical conversion: (6960, 4)
Shape of validation labels after categorical conversion: (1500, 4)
Shape of test labels after categorical conversion: (1772, 4)
Shape of training set after adding width info: (6960, 22, 250, 1)
Shape of validation set after adding width info: (1500, 22, 250, 1)
Shape of test set after adding width info: (1772, 22, 250, 1)
Shape of training set after dimension reshaping: (6960, 250, 1, 22)
Shape of validation set after dimension reshaping: (1500, 250, 1, 22)
Shape of test set after dimension reshaping: (1772, 250, 1, 22)

```

(iii)(CNN-LSTM) Defining the architecture of the hybrid CNN-LSTM model

```

In [ ]: # Building the CNN model using sequential class
def build_model():
    # models = []

    # Conv. block 1
    In1 = keras.Input(shape=(250,1,22) )
    c1 = Conv2D(filters=30, kernel_size=(11,1), padding='same', activation='s
    p1 = MaxPooling2D(pool_size=(4,1), padding='same')(c1) # Read the keras
    b1 = BatchNormalization()(p1)
    d1 = Dropout(0.5)(b1)

    # Conv. block 2
    c2 = Conv2D(filters=60, kernel_size=(9,1), padding='same', activation='s
    p2 = MaxPooling2D(pool_size=(4,1), padding='same')(c2) # Read the keras
    b2 = BatchNormalization()(p2)
    d2 = Dropout(0.6)(b2)

    # Conv. block 3
    c3 = Conv2D(filters=120, kernel_size=(5,1), padding='same', activation='s
    p3 = MaxPooling2D(pool_size=(4,1), padding='same')(c3) # Read the keras
    b3 = BatchNormalization()(p3)
    d3 = Dropout(0.6)(b3)

    # Conv. block 4
    c4 = Conv2D(filters=240, kernel_size=(3,1), padding='same', activation='s
    p4 = MaxPooling2D(pool_size=(4,1), padding='same')(c4) # Read the keras
    b4 = BatchNormalization()(p4)
    d4 = Dropout(0.6)(b4)

    # FC+LSTM layers
    lstm = Flatten()(d4) # Adding a flattening operation to the output of C
    lstm = Dense((120))(lstm) # FC layer with 100 units
    lstm = Reshape((120,1))(lstm) # Reshape my output of FC layer so that it
    lstm = LSTM(10, dropout=0.6, recurrent_dropout=0.1, input_shape=(120,1),

    # Output layer with Softmax activation
    fc = Dense(4, activation='softmax')(lstm) # Output FC layer with softma

    # Define the final model
    final_model = Model(inputs=In1, outputs=[fc])

    # # Printing the model summary
    final_model.compile(loss='categorical_crossentropy',
                        optimizer='adam',
                        metrics=['accuracy'])
    final_model.summary()
    return final_model

```

(iv)(CNN-LSTM) Defining the hyperparameters of the hybrid CNN-LSTM model

```
In [ ]: # Model parameters
learning_rate = 1e-3
epochs = 100
hybrid_cnn_lstm_optimizer = keras.optimizers.Adam(lr=learning_rate)

/Users/lilyzhou/opt/anaconda3/envs/finalProject147/lib/python3.9/site-packages/keras/optimizers/optimizer_v2/adam.py:114: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
```

(v)(CNN-LSTM) Compiling, training and validating the model

```
In [ ]: # Compiling the model
keras.backend.clear_session()

# Training and validating the model

batch_sizes = [32,64,128]
import matplotlib.pyplot as plt

hybrid_cnn_lstm_model = build_model()
# Compiling the model
hybrid_cnn_lstm_model.compile(loss='categorical_crossentropy',
                              optimizer=hybrid_cnn_lstm_optimizer,
                              metrics=['accuracy'])

# Training and validating the model

callback = keras.callbacks.EarlyStopping(monitor='val_loss', patience=20) #
hybrid_cnn_lstm_model_results = hybrid_cnn_lstm_model.fit(x_train,
                                                           y_train,
                                                           batch_size=32,
                                                           epochs=epochs,
                                                           validation_data=(x_valid, y_valid),
                                                           verbose=True,
                                                           callbacks = [callback])
```

2023-03-18 20:47:10.147668: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 250, 1, 22)]	0
conv2d (Conv2D)	(None, 250, 1, 30)	7290
max_pooling2d (MaxPooling2D)	(None, 63, 1, 30)	0
batch_normalization (Batch Normalization)	(None, 63, 1, 30)	120
dropout (Dropout)	(None, 63, 1, 30)	0
conv2d_1 (Conv2D)	(None, 63, 1, 60)	16260
max_pooling2d_1 (MaxPooling2D)	(None, 16, 1, 60)	0
batch_normalization_1 (Batch Normalization)	(None, 16, 1, 60)	240
dropout_1 (Dropout)	(None, 16, 1, 60)	0
conv2d_2 (Conv2D)	(None, 16, 1, 120)	36120
max_pooling2d_2 (MaxPooling2D)	(None, 4, 1, 120)	0
batch_normalization_2 (Batch Normalization)	(None, 4, 1, 120)	480
dropout_2 (Dropout)	(None, 4, 1, 120)	0
conv2d_3 (Conv2D)	(None, 4, 1, 240)	86640
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 240)	0
batch_normalization_3 (Batch Normalization)	(None, 1, 1, 240)	960
dropout_3 (Dropout)	(None, 1, 1, 240)	0
flatten (Flatten)	(None, 240)	0
dense (Dense)	(None, 120)	28920
reshape (Reshape)	(None, 120, 1)	0
lstm (LSTM)	(None, 10)	480
dense_1 (Dense)	(None, 4)	44

=====
Total params: 177,554
Trainable params: 176,654
Non-trainable params: 900
=====

Epoch 1/100

218/218 [=====] - 30s 126ms/step - loss: 1.3888 - accuracy: 0.2530 - val_loss: 1.3825 - val_accuracy: 0.2587

Epoch 2/100

218/218 [=====] - 26s 122ms/step - loss: 1.3784 - accuracy: 0.2855 - val_loss: 1.3781 - val_accuracy: 0.2933

Epoch 3/100

218/218 [=====] - 26s 121ms/step - loss: 1.3601 - accuracy: 0.3174 - val_loss: 1.3142 - val_accuracy: 0.3740

Epoch 4/100

218/218 [=====] - 27s 122ms/step - loss: 1.3147 - accuracy: 0.3671 - val_loss: 1.3202 - val_accuracy: 0.3513

Epoch 5/100

218/218 [=====] - 27s 122ms/step - loss: 1.2898 - accuracy: 0.3856 - val_loss: 1.2761 - val_accuracy: 0.3987

Epoch 6/100

218/218 [=====] - 28s 128ms/step - loss: 1.2602 - accuracy: 0.4088 - val_loss: 1.2645 - val_accuracy: 0.4027

Epoch 7/100

218/218 [=====] - 28s 127ms/step - loss: 1.2473 - accuracy: 0.4164 - val_loss: 1.2164 - val_accuracy: 0.4293

Epoch 8/100

218/218 [=====] - 28s 129ms/step - loss: 1.2289 - accuracy: 0.4332 - val_loss: 1.1946 - val_accuracy: 0.4140

Epoch 9/100

218/218 [=====] - 27s 126ms/step - loss: 1.2133 - accuracy: 0.4376 - val_loss: 1.1867 - val_accuracy: 0.4393

Epoch 10/100

218/218 [=====] - 28s 129ms/step - loss: 1.1955 - accuracy: 0.4527 - val_loss: 1.1772 - val_accuracy: 0.4187

Epoch 11/100

218/218 [=====] - 25s 115ms/step - loss: 1.1801 - accuracy: 0.4611 - val_loss: 1.1564 - val_accuracy: 0.4753

Epoch 12/100

218/218 [=====] - 24s 112ms/step - loss: 1.1599 - accuracy: 0.4744 - val_loss: 1.1605 - val_accuracy: 0.4460

Epoch 13/100

218/218 [=====] - 24s 111ms/step - loss: 1.1481 - accuracy: 0.4839 - val_loss: 1.1220 - val_accuracy: 0.4713

Epoch 14/100

218/218 [=====] - 28s 128ms/step - loss: 1.1344 - accuracy: 0.4907 - val_loss: 1.1273 - val_accuracy: 0.4793

Epoch 15/100

218/218 [=====] - 26s 117ms/step - loss: 1.1247 - accuracy: 0.4894 - val_loss: 1.1509 - val_accuracy: 0.4580

Epoch 16/100

218/218 [=====] - 27s 123ms/step - loss: 1.1074 - accuracy: 0.4981 - val_loss: 1.1036 - val_accuracy: 0.4813

Epoch 17/100

218/218 [=====] - 24s 112ms/step - loss: 1.0924 - accuracy: 0.5076 - val_loss: 1.1046 - val_accuracy: 0.4793

Epoch 18/100
218/218 [=====] - 26s 118ms/step - loss: 1.0796 - accuracy: 0.5194 - val_loss: 1.1236 - val_accuracy: 0.4827
Epoch 19/100
218/218 [=====] - 26s 118ms/step - loss: 1.0717 - accuracy: 0.5162 - val_loss: 1.0904 - val_accuracy: 0.4873
Epoch 20/100
218/218 [=====] - 26s 118ms/step - loss: 1.0579 - accuracy: 0.5283 - val_loss: 1.1677 - val_accuracy: 0.4513
Epoch 21/100
218/218 [=====] - 26s 119ms/step - loss: 1.0630 - accuracy: 0.5341 - val_loss: 1.0984 - val_accuracy: 0.4867
Epoch 22/100
218/218 [=====] - 24s 112ms/step - loss: 1.0345 - accuracy: 0.5434 - val_loss: 1.1001 - val_accuracy: 0.4653
Epoch 23/100
218/218 [=====] - 25s 113ms/step - loss: 1.0343 - accuracy: 0.5447 - val_loss: 1.0842 - val_accuracy: 0.5007
Epoch 24/100
218/218 [=====] - 25s 116ms/step - loss: 1.0207 - accuracy: 0.5468 - val_loss: 1.1349 - val_accuracy: 0.4713
Epoch 25/100
218/218 [=====] - 25s 116ms/step - loss: 1.0206 - accuracy: 0.5563 - val_loss: 1.1191 - val_accuracy: 0.4733
Epoch 26/100
218/218 [=====] - 27s 124ms/step - loss: 1.0016 - accuracy: 0.5619 - val_loss: 1.0863 - val_accuracy: 0.4967
Epoch 27/100
218/218 [=====] - 25s 117ms/step - loss: 1.0061 - accuracy: 0.5585 - val_loss: 1.0849 - val_accuracy: 0.4847
Epoch 28/100
218/218 [=====] - 25s 117ms/step - loss: 0.9898 - accuracy: 0.5644 - val_loss: 1.0872 - val_accuracy: 0.4927
Epoch 29/100
218/218 [=====] - 25s 116ms/step - loss: 0.9812 - accuracy: 0.5615 - val_loss: 1.1203 - val_accuracy: 0.4833
Epoch 30/100
218/218 [=====] - 26s 119ms/step - loss: 0.9754 - accuracy: 0.5721 - val_loss: 1.0724 - val_accuracy: 0.5260
Epoch 31/100
218/218 [=====] - 26s 119ms/step - loss: 0.9756 - accuracy: 0.5720 - val_loss: 1.1420 - val_accuracy: 0.4860
Epoch 32/100
218/218 [=====] - 26s 118ms/step - loss: 0.9658 - accuracy: 0.5843 - val_loss: 1.0819 - val_accuracy: 0.4993
Epoch 33/100
218/218 [=====] - 29s 135ms/step - loss: 0.9469 - accuracy: 0.5874 - val_loss: 1.0912 - val_accuracy: 0.5160
Epoch 34/100
218/218 [=====] - 39s 178ms/step - loss: 0.9495 - accuracy: 0.5861 - val_loss: 1.0730 - val_accuracy: 0.5173
Epoch 35/100
218/218 [=====] - 28s 126ms/step - loss: 0.9444 - accuracy: 0.5948 - val_loss: 1.1131 - val_accuracy: 0.5193
Epoch 36/100
218/218 [=====] - 27s 122ms/step - loss: 0.9452 -

accuracy: 0.5922 - val_loss: 1.0920 - val_accuracy: 0.5173
Epoch 37/100
218/218 [=====] - 25s 114ms/step - loss: 0.9203 -
accuracy: 0.6010 - val_loss: 1.0875 - val_accuracy: 0.5147
Epoch 38/100
218/218 [=====] - 28s 130ms/step - loss: 0.9207 -
accuracy: 0.6080 - val_loss: 1.1145 - val_accuracy: 0.5273
Epoch 39/100
218/218 [=====] - 33s 151ms/step - loss: 0.9098 -
accuracy: 0.6092 - val_loss: 1.1204 - val_accuracy: 0.5133
Epoch 40/100
218/218 [=====] - 37s 172ms/step - loss: 0.9177 -
accuracy: 0.6043 - val_loss: 1.1075 - val_accuracy: 0.5213
Epoch 41/100
218/218 [=====] - 39s 178ms/step - loss: 0.9066 -
accuracy: 0.6154 - val_loss: 1.0882 - val_accuracy: 0.5380
Epoch 42/100
218/218 [=====] - 39s 180ms/step - loss: 0.8933 -
accuracy: 0.6157 - val_loss: 1.0898 - val_accuracy: 0.5267
Epoch 43/100
218/218 [=====] - 43s 195ms/step - loss: 0.8898 -
accuracy: 0.6261 - val_loss: 1.1282 - val_accuracy: 0.5333
Epoch 44/100
218/218 [=====] - 36s 167ms/step - loss: 0.8853 -
accuracy: 0.6177 - val_loss: 1.1274 - val_accuracy: 0.5453
Epoch 45/100
218/218 [=====] - 26s 121ms/step - loss: 0.8816 -
accuracy: 0.6224 - val_loss: 1.1006 - val_accuracy: 0.5440
Epoch 46/100
218/218 [=====] - 28s 128ms/step - loss: 0.8881 -
accuracy: 0.6332 - val_loss: 1.0269 - val_accuracy: 0.6000
Epoch 47/100
218/218 [=====] - 34s 157ms/step - loss: 0.8632 -
accuracy: 0.6425 - val_loss: 1.1025 - val_accuracy: 0.5507
Epoch 48/100
218/218 [=====] - 46s 213ms/step - loss: 0.8664 -
accuracy: 0.6392 - val_loss: 1.0858 - val_accuracy: 0.5747
Epoch 49/100
218/218 [=====] - 44s 200ms/step - loss: 0.8622 -
accuracy: 0.6404 - val_loss: 1.0234 - val_accuracy: 0.5833
Epoch 50/100
218/218 [=====] - 42s 194ms/step - loss: 0.8523 -
accuracy: 0.6526 - val_loss: 1.0947 - val_accuracy: 0.5593
Epoch 51/100
218/218 [=====] - 40s 183ms/step - loss: 0.8420 -
accuracy: 0.6557 - val_loss: 1.0424 - val_accuracy: 0.5860
Epoch 52/100
218/218 [=====] - 43s 196ms/step - loss: 0.8377 -
accuracy: 0.6582 - val_loss: 1.0468 - val_accuracy: 0.5927
Epoch 53/100
218/218 [=====] - 43s 196ms/step - loss: 0.8355 -
accuracy: 0.6658 - val_loss: 1.1237 - val_accuracy: 0.5513
Epoch 54/100
218/218 [=====] - 41s 189ms/step - loss: 0.8258 -
accuracy: 0.6690 - val_loss: 1.0386 - val_accuracy: 0.6073
Epoch 55/100

218/218 [=====] - 41s 188ms/step - loss: 0.8243 - accuracy: 0.6658 - val_loss: 1.0317 - val_accuracy: 0.6053
Epoch 56/100
218/218 [=====] - 38s 173ms/step - loss: 0.8162 - accuracy: 0.6700 - val_loss: 1.0584 - val_accuracy: 0.5920
Epoch 57/100
218/218 [=====] - 38s 172ms/step - loss: 0.8020 - accuracy: 0.6823 - val_loss: 1.0689 - val_accuracy: 0.6193
Epoch 58/100
218/218 [=====] - 38s 173ms/step - loss: 0.8004 - accuracy: 0.6826 - val_loss: 1.0095 - val_accuracy: 0.6267
Epoch 59/100
218/218 [=====] - 36s 166ms/step - loss: 0.7959 - accuracy: 0.6931 - val_loss: 1.0293 - val_accuracy: 0.6320
Epoch 60/100
218/218 [=====] - 45s 204ms/step - loss: 0.8029 - accuracy: 0.6886 - val_loss: 0.9983 - val_accuracy: 0.6167
Epoch 61/100
218/218 [=====] - 41s 190ms/step - loss: 0.7871 - accuracy: 0.6935 - val_loss: 0.9741 - val_accuracy: 0.6433
Epoch 62/100
218/218 [=====] - 40s 183ms/step - loss: 0.7792 - accuracy: 0.6934 - val_loss: 1.0127 - val_accuracy: 0.6187
Epoch 63/100
218/218 [=====] - 41s 190ms/step - loss: 0.7799 - accuracy: 0.6950 - val_loss: 0.9592 - val_accuracy: 0.6427
Epoch 64/100
218/218 [=====] - 41s 189ms/step - loss: 0.7686 - accuracy: 0.6986 - val_loss: 0.9687 - val_accuracy: 0.6373
Epoch 65/100
218/218 [=====] - 41s 189ms/step - loss: 0.7457 - accuracy: 0.7108 - val_loss: 0.9956 - val_accuracy: 0.6433
Epoch 66/100
218/218 [=====] - 42s 194ms/step - loss: 0.7464 - accuracy: 0.7151 - val_loss: 1.0392 - val_accuracy: 0.6293
Epoch 67/100
218/218 [=====] - 42s 193ms/step - loss: 0.7421 - accuracy: 0.7152 - val_loss: 1.0770 - val_accuracy: 0.6227
Epoch 68/100
218/218 [=====] - 43s 197ms/step - loss: 0.7391 - accuracy: 0.7131 - val_loss: 1.0092 - val_accuracy: 0.6207
Epoch 69/100
218/218 [=====] - 40s 185ms/step - loss: 0.7318 - accuracy: 0.7220 - val_loss: 1.0386 - val_accuracy: 0.6273
Epoch 70/100
218/218 [=====] - 37s 172ms/step - loss: 0.7245 - accuracy: 0.7266 - val_loss: 0.9895 - val_accuracy: 0.6320
Epoch 71/100
218/218 [=====] - 24s 111ms/step - loss: 0.7143 - accuracy: 0.7325 - val_loss: 0.9589 - val_accuracy: 0.6407
Epoch 72/100
218/218 [=====] - 25s 117ms/step - loss: 0.7122 - accuracy: 0.7282 - val_loss: 0.9545 - val_accuracy: 0.6407
Epoch 73/100
218/218 [=====] - 25s 117ms/step - loss: 0.7008 - accuracy: 0.7346 - val_loss: 0.9777 - val_accuracy: 0.6527

Epoch 74/100
218/218 [=====] - 25s 112ms/step - loss: 0.6979 - accuracy: 0.7362 - val_loss: 0.9727 - val_accuracy: 0.6480
Epoch 75/100
218/218 [=====] - 24s 111ms/step - loss: 0.6917 - accuracy: 0.7362 - val_loss: 0.9871 - val_accuracy: 0.6313
Epoch 76/100
218/218 [=====] - 22s 99ms/step - loss: 0.6941 - accuracy: 0.7371 - val_loss: 0.9648 - val_accuracy: 0.6360
Epoch 77/100
218/218 [=====] - 19s 85ms/step - loss: 0.7047 - accuracy: 0.7330 - val_loss: 0.9653 - val_accuracy: 0.6480
Epoch 78/100
218/218 [=====] - 18s 84ms/step - loss: 0.6903 - accuracy: 0.7386 - val_loss: 1.0132 - val_accuracy: 0.6167
Epoch 79/100
218/218 [=====] - 18s 84ms/step - loss: 0.6685 - accuracy: 0.7486 - val_loss: 1.0329 - val_accuracy: 0.6287
Epoch 80/100
218/218 [=====] - 18s 85ms/step - loss: 0.6888 - accuracy: 0.7404 - val_loss: 0.9615 - val_accuracy: 0.6353
Epoch 81/100
218/218 [=====] - 19s 85ms/step - loss: 0.6706 - accuracy: 0.7447 - val_loss: 0.9342 - val_accuracy: 0.6507
Epoch 82/100
218/218 [=====] - 18s 84ms/step - loss: 0.6798 - accuracy: 0.7468 - val_loss: 0.9883 - val_accuracy: 0.6340
Epoch 83/100
218/218 [=====] - 18s 83ms/step - loss: 0.6730 - accuracy: 0.7444 - val_loss: 0.9681 - val_accuracy: 0.6520
Epoch 84/100
218/218 [=====] - 19s 85ms/step - loss: 0.6567 - accuracy: 0.7540 - val_loss: 0.9574 - val_accuracy: 0.6513
Epoch 85/100
218/218 [=====] - 18s 84ms/step - loss: 0.6687 - accuracy: 0.7473 - val_loss: 0.9385 - val_accuracy: 0.6567
Epoch 86/100
218/218 [=====] - 18s 84ms/step - loss: 0.6514 - accuracy: 0.7547 - val_loss: 0.9621 - val_accuracy: 0.6527
Epoch 87/100
218/218 [=====] - 16s 71ms/step - loss: 0.6498 - accuracy: 0.7524 - val_loss: 0.9330 - val_accuracy: 0.6667
Epoch 88/100
218/218 [=====] - 14s 66ms/step - loss: 0.6520 - accuracy: 0.7547 - val_loss: 0.9660 - val_accuracy: 0.6340
Epoch 89/100
218/218 [=====] - 15s 67ms/step - loss: 0.6533 - accuracy: 0.7517 - val_loss: 0.9687 - val_accuracy: 0.6473
Epoch 90/100
218/218 [=====] - 15s 66ms/step - loss: 0.6489 - accuracy: 0.7547 - val_loss: 0.9672 - val_accuracy: 0.6547
Epoch 91/100
218/218 [=====] - 15s 67ms/step - loss: 0.6363 - accuracy: 0.7588 - val_loss: 0.9459 - val_accuracy: 0.6413
Epoch 92/100
218/218 [=====] - 15s 67ms/step - loss: 0.6472 - a

```

ccuracy: 0.7585 - val_loss: 0.8979 - val_accuracy: 0.6673
Epoch 93/100
218/218 [=====] - 15s 69ms/step - loss: 0.6253 - a
ccuracy: 0.7611 - val_loss: 0.8922 - val_accuracy: 0.6793
Epoch 94/100
218/218 [=====] - 14s 66ms/step - loss: 0.6239 - a
ccuracy: 0.7651 - val_loss: 0.9626 - val_accuracy: 0.6453
Epoch 95/100
218/218 [=====] - 13s 62ms/step - loss: 0.6283 - a
ccuracy: 0.7675 - val_loss: 0.9193 - val_accuracy: 0.6667
Epoch 96/100
218/218 [=====] - 13s 61ms/step - loss: 0.6360 - a
ccuracy: 0.7573 - val_loss: 0.9163 - val_accuracy: 0.6633
Epoch 97/100
218/218 [=====] - 13s 61ms/step - loss: 0.6319 - a
ccuracy: 0.7651 - val_loss: 0.9545 - val_accuracy: 0.6380
Epoch 98/100
218/218 [=====] - 13s 61ms/step - loss: 0.6256 - a
ccuracy: 0.7635 - val_loss: 0.8699 - val_accuracy: 0.6860
Epoch 99/100
218/218 [=====] - 13s 62ms/step - loss: 0.6233 - a
ccuracy: 0.7670 - val_loss: 0.8874 - val_accuracy: 0.6880
Epoch 100/100
218/218 [=====] - 14s 64ms/step - loss: 0.6180 - a
ccuracy: 0.7684 - val_loss: 0.8565 - val_accuracy: 0.6860

```

(vi)(CNN-LSTM) Visualizing the accuracy and loss trajectory

```


In [ ]: import matplotlib.pyplot as plt

# Plotting accuracy trajectory
plt.plot(hybrid_cnn_lstm_model_results.history['accuracy'])
plt.plot(hybrid_cnn_lstm_model_results.history['val_accuracy'])
plt.title('Hybrid CNN-LSTM model accuracy trajectory')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()


# Plotting loss trajectory
plt.plot(hybrid_cnn_lstm_model_results.history['loss'], 'o')
plt.plot(hybrid_cnn_lstm_model_results.history['val_loss'], 'o')
plt.title('Hybrid CNN-LSTM model loss trajectory')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()

```

Hybrid CNN-LSTM model accuracy trajectory



Hybrid CNN-LSTM model loss trajectory



(vii)(CNN-LSTM) Testing the performance of the hybrid CNN-LSTM model on the held out test set

```
In [ ]: ## Testing the hybrid CNN-LSTM model

hybrid_cnn_lstm_score = hybrid_cnn_lstm_model.evaluate(x_test, y_test, verbose=0)
print('Test accuracy of the hybrid CNN-LSTM model:', hybrid_cnn_lstm_score[1])

Test accuracy of the hybrid CNN-LSTM model: 0.6834085583686829
```

```
In [ ]:
```